HackHunt Crew

Shared memory, bounded context.

ContextCon | Theme 08: Do More With Fewer Tokens

Four roles. One Meko datapack. Measured client token reduction.

The context-history problem

Stop copying every earlier fact into every prompt.

Scout extracts facts. Verifier checks seven rules.

RedTeam flags hard violations. Builder proposes a project.

Baseline: all accumulated facts. Meko: up to five scoped memories.

Meko at the center

Shared facts with source boundaries.

Candidate-specific memory conversation scopes shared across roles.

Owner-scoped role traces and client token attribution.

Criteria uploaded to knowledge base; complete reports stored as artifacts.

Current supplied evidence takes priority over remembered claims.

Original experiment failed quality

We kept the failure, then changed the decision policy.

Development v3: 86.2% fewer client tokens, but Meko 65/72 correct versus baseline 69/72. Not a same-quality result.

Added a label-blind gate over extracted facts: unknown is not failure.

Post-hoc development rescoring is separate from prospective testing.

Prospective held-out test

74.0% fewer client LLM tokens.

12 new synthetic cases, 3 sequential passes, fixed model and prompts.

Baseline: 411,058 tokens, 35/36 correct, 120 successful calls.

Meko: 106,764 tokens, 36/36 correct, 120 successful calls.

No event names or ground-truth labels used by the gate.

Per-pass results

Full history grows. Scoped context stays bounded.

Pass	1	2	3
Baseline	61,483	136,890	212,685
Meko	31,891	38,250	36,623

Mean tokens: baseline 137,019.3; Meko 35,588.0.

Limits and reproduction

Small synthetic test, explicit measurement boundaries.

Sequential passes, not independent statistical trials.

Naive full-history baseline, not optimized local retrieval.

Client tokens exclude Meko backend inference, retries and interrupted attempts.

No general production or whole-system cost claim.

Code, reports, frozen plan and 13 unit tests are public.